

# **Operating Systems**

## **Lecture – 8-10**

### **CPU Scheduling Algorithms**

**Dr. Shamim Akhter**

# Problem



- You are the cook at the State St. Diner
  - Customers continually enter and place their orders
  - Dishes take varying amounts of time to prepare
- What is your goal?
- Which strategy achieves this goal?

# Correspondence **DINER** - Multitasking Operating System

- Ordering Manager  $\leftrightarrow$  CPU Scheduler
- Cook  $\leftrightarrow$  CPU
- Client  $\leftrightarrow$  User
- Order  $\leftrightarrow$  Process
- Dish category  $\leftrightarrow$  CPU/I-O Burst

# CPU/I-O Burst

- Process alternates between CPU and I/O bursts
  - **CPU-bound** jobs: Long CPU bursts



- **I/O-bound**: Short CPU bursts

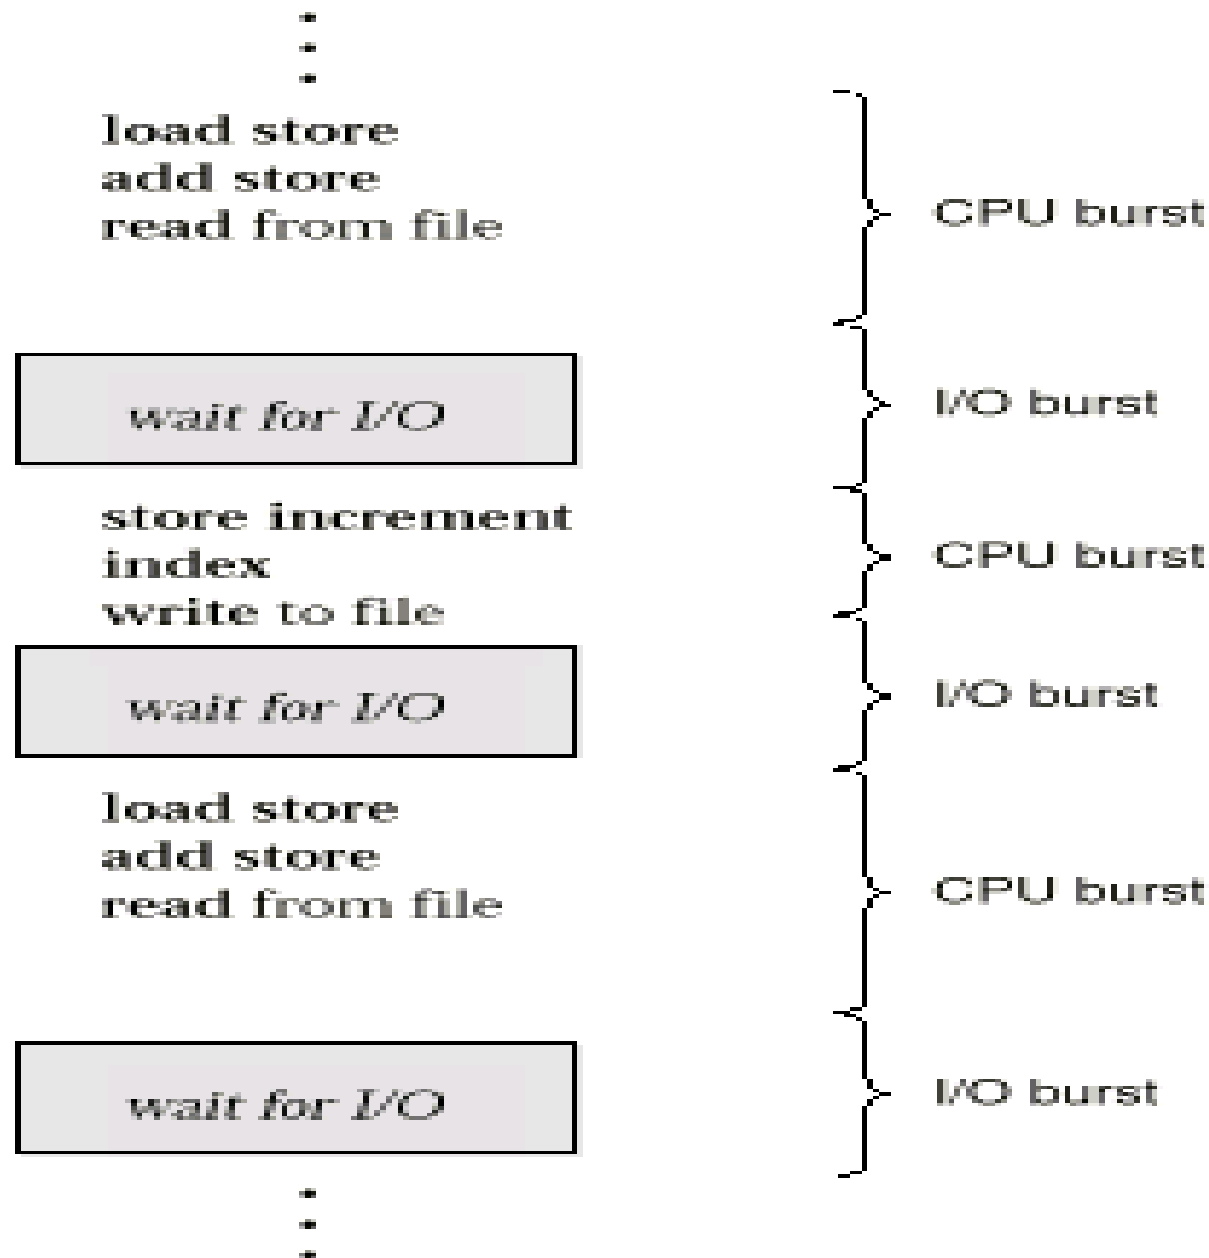


- **I/O burst** = process idle, switch to another "for free"

# CPU – I/O Burst Cycle

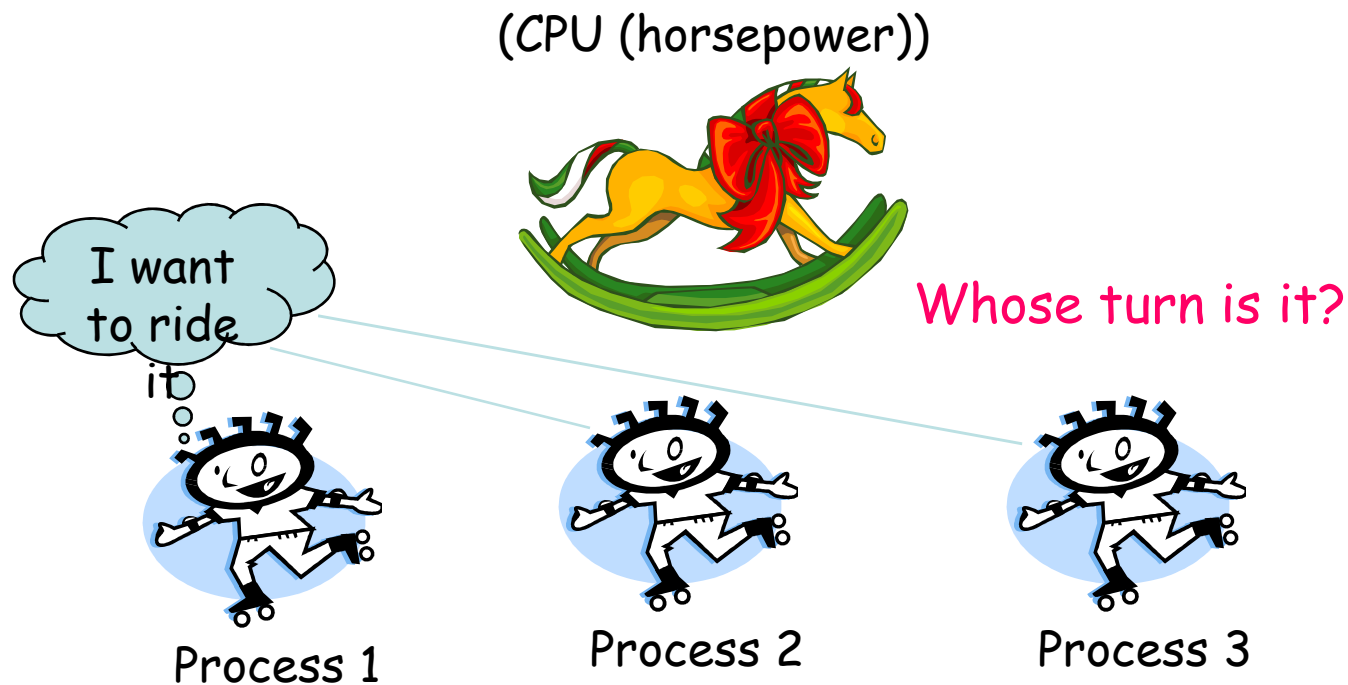
- Process execution consists of a cycle of CPU execution and I/O wait
  - Processes move back & forth between these two states
- A process execution begins with a CPU burst, followed by an I/O burst and then another CPU burst and so on
- An alternating sequence of CPU & I/O bursts is shown on next slide

# Alternating Sequence of CPU and I/O Bursts



# Scheduling

- Deciding which process/thread should occupy a resource (CPU, disk, etc.)



# CPU Scheduler

- When CPU becomes idle
  - OS must select one of the processes in the Ready Queue to be executed.
- **CPU scheduler**
  - Is the OS code that implements the CPU scheduling algorithm .
  - Selects a process from the processes in memory that are ready to execute (from Ready Queue), and allocates the CPU to one of them.
- CPU scheduling decisions may take place when a process:
  1. Switches from running to waiting state
  2. Switches from running to ready state (e.g. when time slice of a process expires or an interrupt occurs)
  3. Switches from waiting to ready state. (e.g. on completion of I/O)
  4. On arrival of a new process
  5. Terminates



# Dispatcher

- An important component involved in CPU scheduling
- The OS code that
  - takes the CPU away from the current process and hands it over to the newly scheduled process .
- Dispatcher gives control of the CPU to the process selected by the short-term scheduler
- Dispatcher performs following functions:
  - switching context
  - switching to user mode
  - jumping to the proper location in the user program to restart that program
- *Dispatch latency* – time it takes for the dispatcher to stop one process and start executing another

# Preemptive vs Non Preemptive Process

## Non-preemptive

- ✓ Process runs until voluntarily relinquish CPU
  - process blocks on an event (e.g., I/O or synchronization)
  - process terminates
  - process periodically calls the `yield()` system call to give up the CPU
- ✓ Only suitable for domains where processes can be trusted to relinquish the CPU

## Preemptive

- The scheduler actively interrupts and reschedules an executing process
- Required when applications cannot be trusted to yield
- Incurs some overhead

# Scheduling Criteria

- **CPU utilization** – percentage of time the CPU is not idle.
- **Throughput** – completed processes per unit of time
- **Waiting time** – time spent in the ready queue
  - For Non preemptive Algos =  $\text{Starting.Time} - \text{Arrival.Time}$
  - For Preemptive Algos =  $\text{Finish.Time} - \text{Arrival.Time} - \text{Burst.Time}$
- **Turnaround time** – amount of time to execute a particular process.
  - $\text{Finish.Time} - \text{Arrival.Time}$
- **Response time** – response latency,
  - amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

# The Perfect Scheduler

- Max CPU utilization= keep all devices busy
- Max throughput =Maximize jobs/time
- Min turnaround time
- Min waiting time
- Min response time =latency, first response
- Fairness: everyone makes progress, no one starves

# Single CPU–Scheduling Algorithms

- First Come First Serve (FCFS)
- Smallest Next CPU Burst
  - Shortest Job First (SJF)
  - Shortest Remaining Time First (SRTF)
- Round Robin
- Priority Scheduling

# First Come First Serve

- Simplest CPU scheduling algorithm
- Non preemptive
- The process that requests the CPU first is allocated the CPU first
- Implemented with a FIFO queue.
  - when a process enters the Ready Queue, its PCB is linked on to the tail of the Queue. When the CPU is free it is allocated to the process at the head of the Queue.
- Limitations
  - FCFS **favor long processes as compared to short ones**. (Convoy effect)
  - **Average waiting time depends on arrival order**
  - Average waiting time is often quite long
  - FCFS is non-preemptive, so it is trouble some for time sharing systems

# Convoy Effect

“A convoy effect happens when a set of processes need to use a resource for a short time, and one process holds the resource for a long time, blocking all of the other processes. Causes poor utilization of the other resources in the system”

# FCFS – Example

The process that enters the ready queue first is scheduled first, regardless of the size of its next CPU burst

Example:	Process	Burst Time
	P1	24
	P2	3
	P3	3

Suppose that processes arrive into the system in the order:  
P1, P2 , P3



# FCFS – Example

Processes are served in the order: P1, P2, P3

The **Gantt Chart** for the schedule is:



Waiting times P1 = 0; P2 = 24; P3 = 27

Average waiting time:  $(0+24+27)/3 = 17$

# FCFS – Example

Suppose that processes arrive in the order: P2 , P3 , P1 .  
The Gantt chart for the schedule is:



Waiting time for P1 = 6; P2 = 0; P3 = 3

Average waiting time:  $(6 + 0 + 3)/3 = 3$

# FCFS - Example

Process	Duration/B.T	Order	Arrival Time
P1	24	1	0
P2	3	2	3
P3	4	3	4

The final schedule:



P1 waiting time: 0-0

P2 waiting time: 24-3

P3 waiting time: 27-4

The average waiting time:  
 $(0+21+23)/3 = 14.667$

# FCFS – Example

- Draw the graph (Gantt chart) and compute average waiting time for the following processes using **FCFS** Scheduling algorithm.

<u>Process</u>	<u>Arrival time</u>	<u>Burst Time</u>
P1	1	16
P2	5	3
P3	6	4
P4	9	2

# SJF & SRTF Scheduling...

- When the CPU is available it is assigned to the process that has the smallest next CPU burst.
- If two processes have the same length next CPU bursts, FCFS scheduling is used to break the tie.

## Comes in two flavors

- **Shortest Job First (SJF)**
  - It's a non preemptive algorithm.
- **Shortest Remaining Time First (SRTF)**
  - It's a Preemptive algorithm.

# SJF Example

Process	Duration/B.T	Order	Arrival Time
P1	6	1	0
P2	8	2	0
P3	7	3	0
P4	3	4	0



P4 waiting time: 0-0  
P1 waiting time: 3-0  
P3 waiting time: 9-0  
P2 waiting time: 16-0

The total running time is: 28  
The average waiting time (AWT):  
 $(0+3+9+16)/4 = 7$  time units

# SRTF Example

Process	Duration	Order	Arrival Time
P1	10	1	0
P2	2	2	2



The average waiting time (AWT):  
P1 waiting time:  $4 - 2 = 2$        $(0 + 2) / 2 = 1$   
P2 waiting time: 0

Now run this using SJF!

# SJF & SRTF – Example

Draw the graph (Gantt chart) and compute waiting time and turn around time for the following processes using **SJF** & **SRTF** Scheduling algorithm.

For SJF consider all processes arrive at time 0 in sequence P1, P2, P3, P4.

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P1	0	8
P2	1	4
P3	2	9
P4	3	5



# SJF & SRTF – Example

Draw the graph (Gantt chart) and compute waiting time and turn around time for the following processes using **SJF** & **SRTF** Scheduling algorithm.

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P1	0	5
P2	1	2
P3	2	3
P4	3	1

# SJF & SRTF – Example

Draw the graph (Gantt chart) and compute waiting time and turn around time for the following processes using **SJF** & **SRTF** Scheduling algorithm.

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P1	0	9
P2	3	6
P3	6	2
P4	9	1

# SJF & SRTF Scheduling

## \$100 QUESTION

How to compute the next CPU burst?



- This algorithm cannot be implemented at the level of short-term CPU scheduling. There is **no way to know the length of the next CPU burst**.
- One approach is to try to approximate.
- We may not *know* the length of the next CPU burst, but we may be able to **predict its value**.
- We expect that the next CPU burst will be similar in length to the previous ones.
  - Thus, by computing an approximation of the length of the next CPU burst, we can pick the process with the shortest predicted CPU burst.

# SJF & SRTF Scheduling...

## Exponential Averaging

Estimation based on historical data

$t_n$  = Actual length of  $n^{th}$  CPU burst

$\tau_n$  = Estimate for  $n^{th}$  CPU burst

$\tau_{n+1}$  = Estimate for  $n+1^{st}$  CPU burst

$\alpha, 0 \leq \alpha \leq 1$

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$

# Exponential Averaging...

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$

- Plugging in value for  $\tau_n$ , we get

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) [\alpha t_{n-1} + (1 - \alpha) \tau_{n-1}]$$

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \alpha t_{n-1} + (1 - \alpha)^2 \tau_{n-1}$$

- Again plugging in value for  $\tau_{n-1}$ , we get

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \alpha t_{n-1} + (1 - \alpha)^2 [\alpha t_{n-2} + (1 - \alpha) \tau_{n-2}]$$

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \alpha t_{n-1} + (1 - \alpha)^2 \alpha t_{n-2} + (1 - \alpha)^3 \tau_{n-2}$$

- Continuing like this results in

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots + (1 - \alpha)^j \alpha t_{n-j} + \dots + (1 - \alpha)^{n+1} \tau_0$$

# Exponential Averaging...

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$

Lets take two extreme values of  $\alpha$

**If  $\alpha = 0$**

$$\tau_{n+1} = \tau_n$$

Next CPU burst estimate will be exactly equal to previous CPU burst estimate.

**If  $\alpha = 1$**

$$\tau_{n+1} = t_n$$

Next CPU burst estimate will be equal to previous actual CPU burst.

# Exponential Averaging...


$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots + (1 - \alpha)^j \alpha t_{n-j} + \dots + (1 - \alpha)^{n+1} \tau_0$$

Typical value used for  $\alpha$  is  $\frac{1}{2}$ . With this value, our  $(n+1)^{\text{st}}$  estimate is

$$\tau_{n+1} = t_n/2 + t_{n-1}/2^2 + t_{n-2}/2^3 + t_{n-3}/2^4 + \dots$$

Note that as we move back in time we are giving less weight-age to the previous CPU bursts. Older histories are given exponentially less weight-age

# Prediction of the Length of the Next CPU Burst



$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$

$$\tau_0 = 10$$

$$\tau_1 = .5 * 6 + .5 * 10 = 8$$

$$\tau_2 = .5 * 4 + .5 * 8 = 6$$

$$\tau_3 = .5 * 6 + .5 * 6 = 6$$

$$\tau_4 = .5 * 4 + .5 * 6 = 5$$

$$\tau_5 = .5 * 13 + .5 * 5 = 9$$

...

...

CPU burst ( $t_i$ )	6	4	6	4	13	13	13	...
---------------------	---	---	---	---	----	----	----	-----

"guess" ( $\tau_i$ )	10	8	6	6	5	9	11	12	...
----------------------	----	---	---	---	---	---	----	----	-----




# Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum**  $q$ ), usually 10-100 milliseconds.
- After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process **gets  $1/n$  of the CPU time in chunks of at most  $q$  time units** at once. No process waits more than  $(n-1)q$  time units.
- Timer interrupts every quantum to schedule next process
- **Performance**
  - $q$  large  $\Rightarrow$  FIFO
  - $q$  small  $\Rightarrow q$  must be large with respect to context switch, otherwise overhead is too high

# Example of RR with Time Quantum = 4


<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- The Gantt chart is:



- Typically, higher average turnaround than SJF, but better *response*
- q should be large compared to context switch time
- q usually 10ms to 100ms

# Time Quantum and Context Switch Time



- One process of 10 time units
  - If  $q$  is 12, finish in less than 1  $q$
  - If  $q$  is 6, requires 2  $q$ , one context switch
  - If  $q$  is 1, requires 10  $q$  and 9 context switches

# Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority  
(smallest integer  $\equiv$  highest priority)
  - Preemptive
  - Nonpreemptive
- SJF is priority scheduling
  - where priority is the inverse of predicted next CPU burst time
  - > Shortest Job, Highest Priority
- Problem  $\equiv$  Starvation – low priority processes may never execute
- Solution  $\equiv$  Aging – as time progresses increase the priority of the process


# Priority Scheduling

- **Priority Scheduling**
  - Choose next job based on priority
  - For SJF, priority = expected CPU burst
  - Can be either preemptive or non-preemptive
- **Priority schedulers can approximate other scheduling algorithms**
  - **$P = \text{arrival time} \Rightarrow \text{FIFO}$**
  - **$P = \text{now} - \text{arrival time} \Rightarrow \text{LIFO}$**
  - **$P = \text{job length} \Rightarrow \text{SJF}$**
- **Problem:**
  - Starvation: jobs can wait indefinitely
- **Solution to starvation**
  - Age processes: increase priority as a function of waiting time

# Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	10	3
$P_2$	1	1
$P_3$	2	4
$P_4$	1	5
$P_5$	5	2


- Priority scheduling Gantt Chart




- Non-Preemptive Waiting time =  $0 + 1 + 6 + 16 + 18 = 41$
- Average waiting time = 8.2 msec

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2


First-Come-First Served (FCFS)




Round-Robin (RR),  $q = 1$



Shortest Process  
Next (SPN)



Shortest Remaining  
Time (SRT)



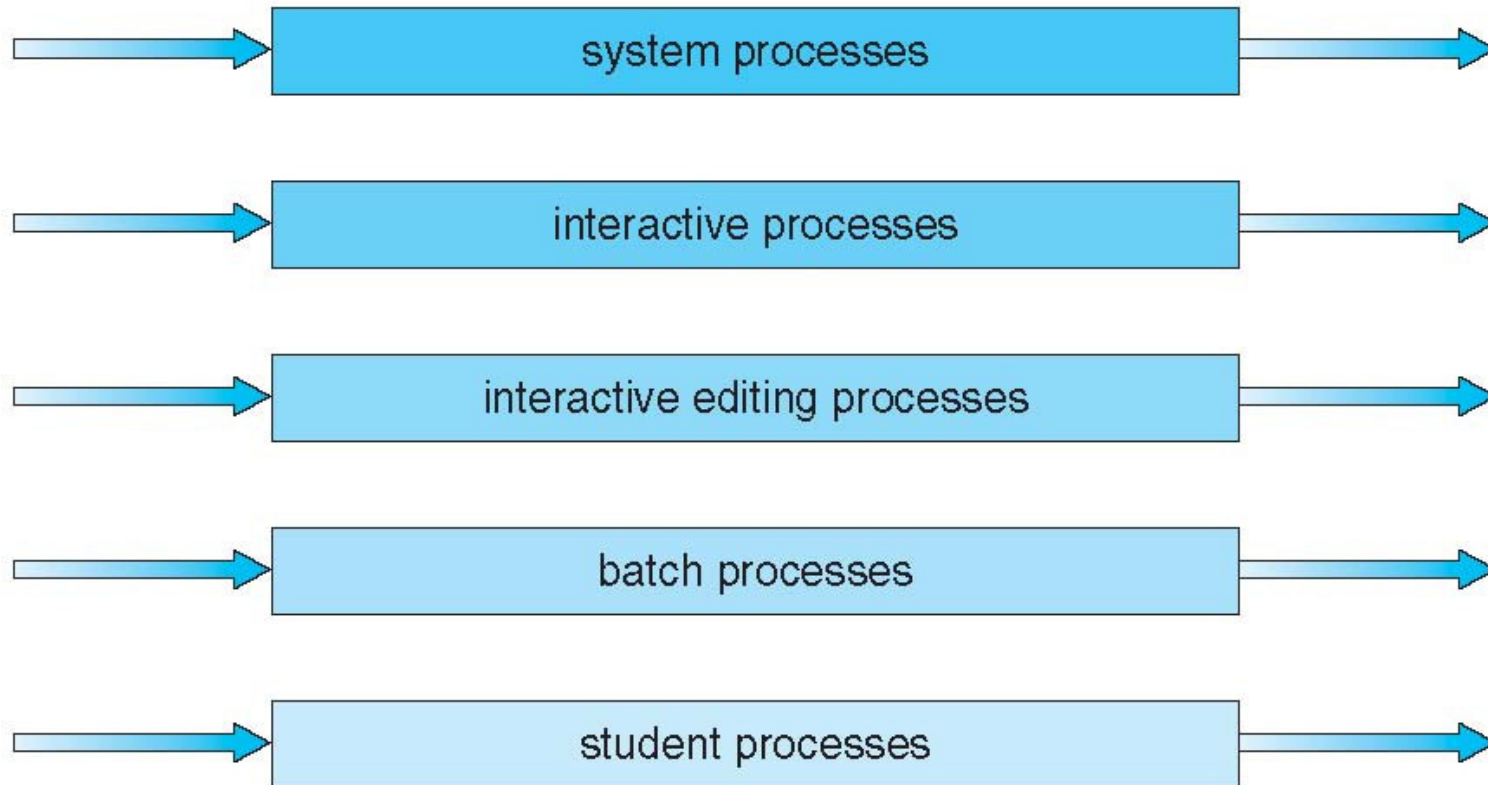


# Multilevel Queue

- Ready queue is partitioned into separate queues, eg:
  - **foreground** (interactive)
  - **background** (batch)
- Process permanently in a given queue
- Each queue has its own scheduling algorithm:
  - foreground – RR
  - background – FCFS
- Scheduling must be done between the queues:
  - Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
  - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes;
  - i.e., 80% to foreground in RR, 20% to background in FCFS

# Multilevel Queue Scheduling

highest priority

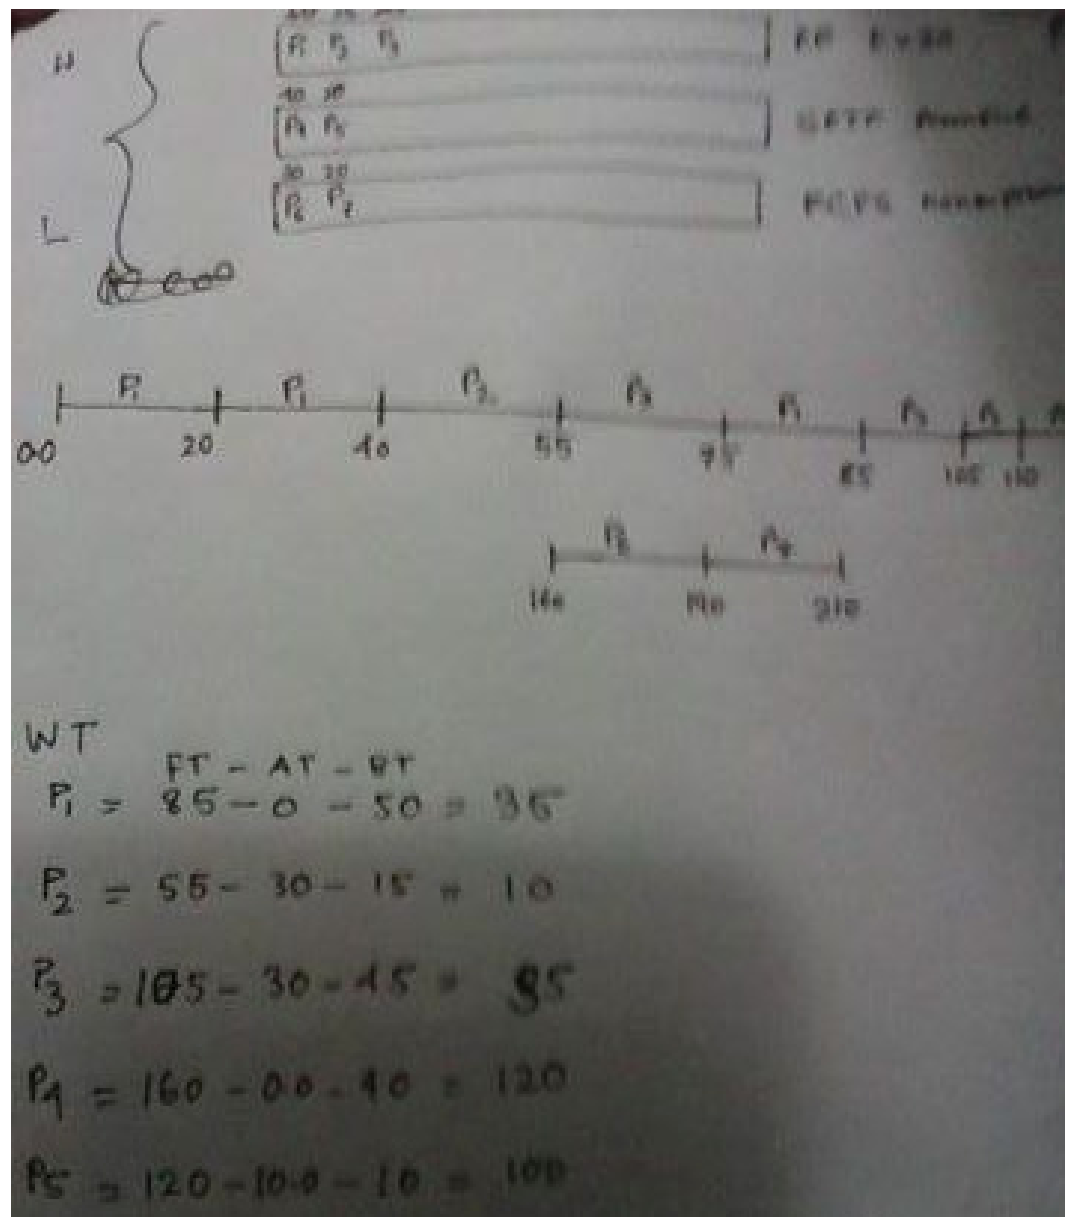


lowest priority

4. By “**Multilevel Queue Fixed priority**” scheduling algorithm, draw the CPU scheduling Gantt chart and complete the table for the following given processes information. (10)

	Queue Name	Process	Burst Time	Arrival Time	Algorithm
Higher Priority	1 <sup>st</sup> Foreground	P1	50	0.0	RR interval: 20 Preemptive
		P2	15	30.0	
		P3	45	30.0	
Lower Priority	2 <sup>nd</sup> Foreground	P4	40	0.0	SRTF Preemptive
		P5	10	10.0	
	Background	P6	30	60	FCFS Non-preemptive
		P7	20	130	


	P1	P2	P3	P4	P5	P6	P7
Waiting time	35	10	35	120	100	100	60
Turnaround time	85	25	80	160	110	130	80



$$P_6 = 190 - 60 - 30 = 100$$

$$P_7 = 210 - 130 - 20 = 60$$

# Multilevel Feedback Queue



# Course Materials

@Galvin-6.1-6.3