


# Lecture 11-13

## Synchronization

---

Dr. Shamim Akhter





# Processes

## ■ A process contains everything needed for execution

- Address space
- PCB

## □ So far, **processes have limited ability** to pass data

- Parents get one chance to pass everything at fork()
- **But what if the child wants to talk back?**
- And, what about processes with different ancestry?

## ■ Yet sometimes processes **may wish to cooperate**

- But how to communicate? Each process is an island
- The OS needs to get involved to bridge the gap
- OS provides system calls to support **Inter-Process Communication (IPC)**





# Interprocess Communication

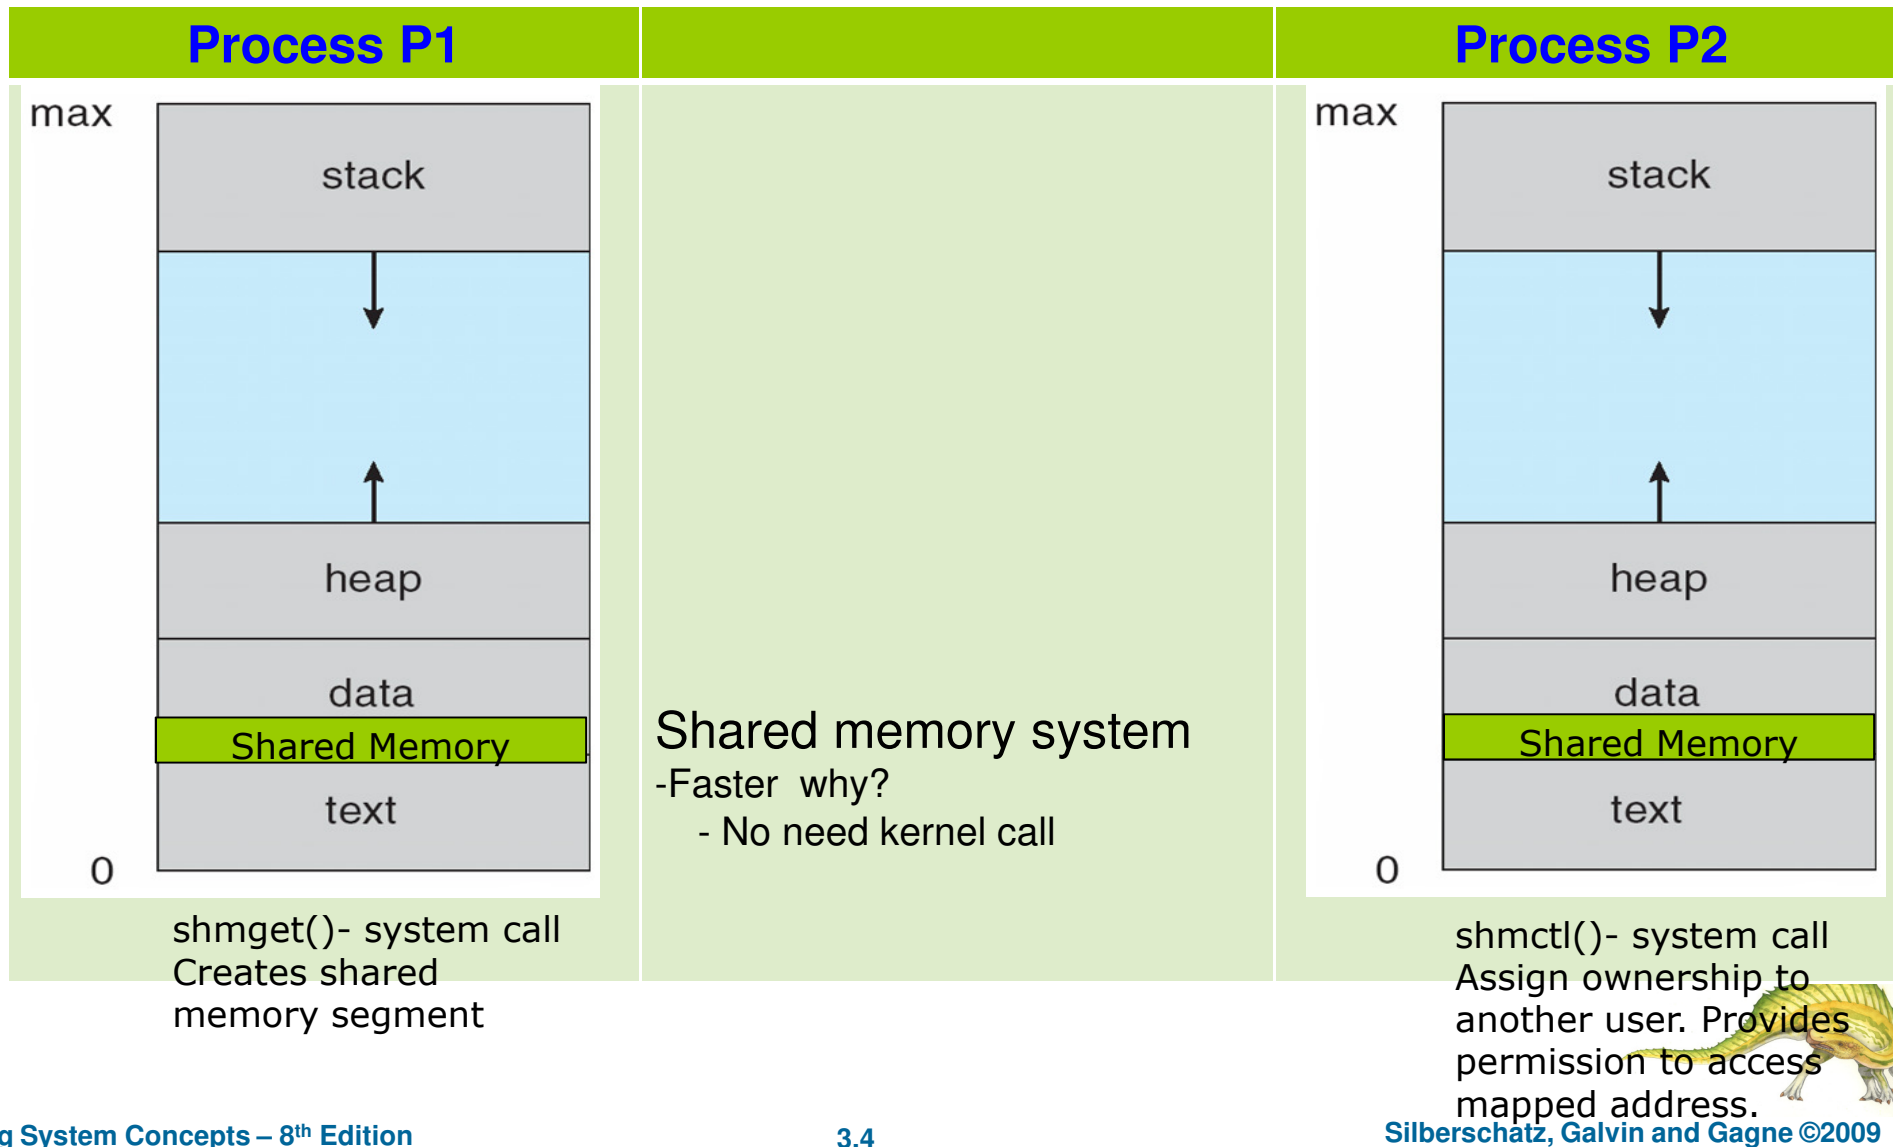
---

- Cooperating processes need **interprocess communication (IPC)**
- Two IPC models -OS provides mechanisms to communicate (System Call)
  - **Shared memory**
    - ▶ multiple processes can read/write same physical portion of memory;
    - ▶ System call to declare shared region
  - **Message passing**
    - ▶ communication channel provided through send()/receive() system calls





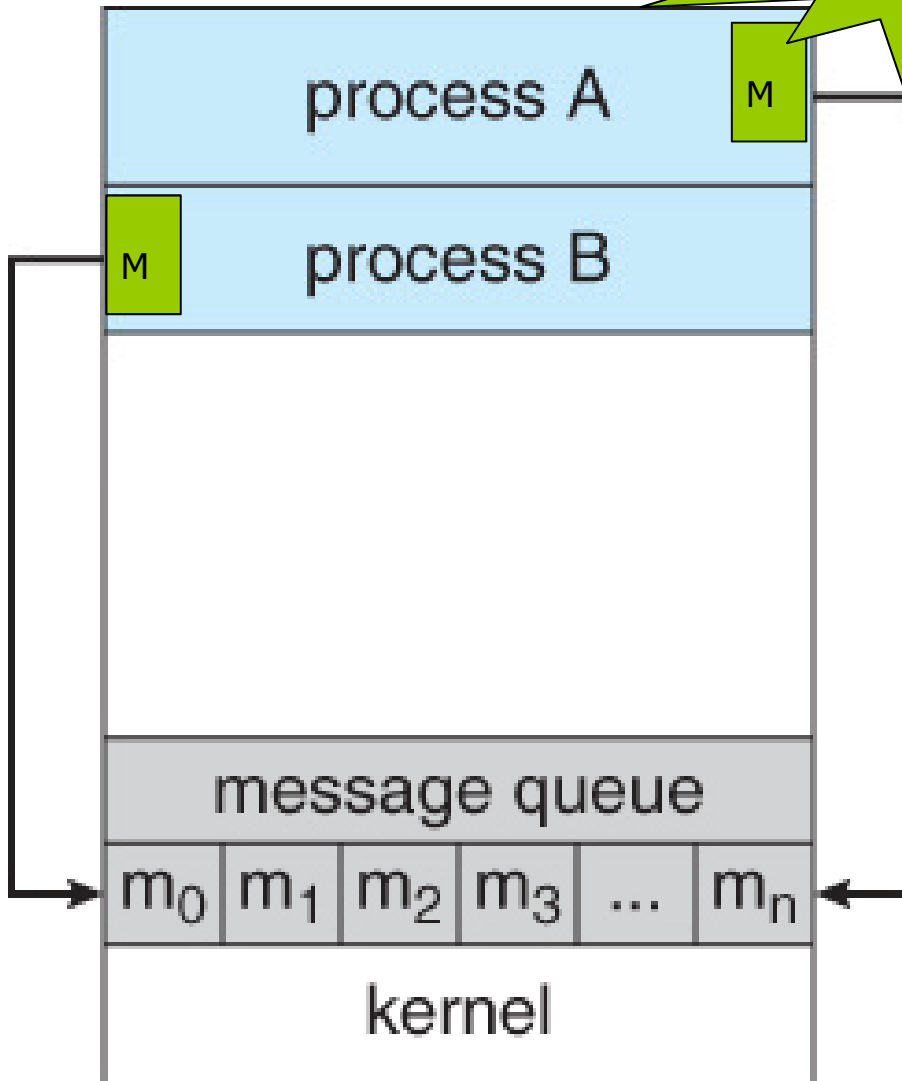
# Shared Memory- 3.3.5(section)



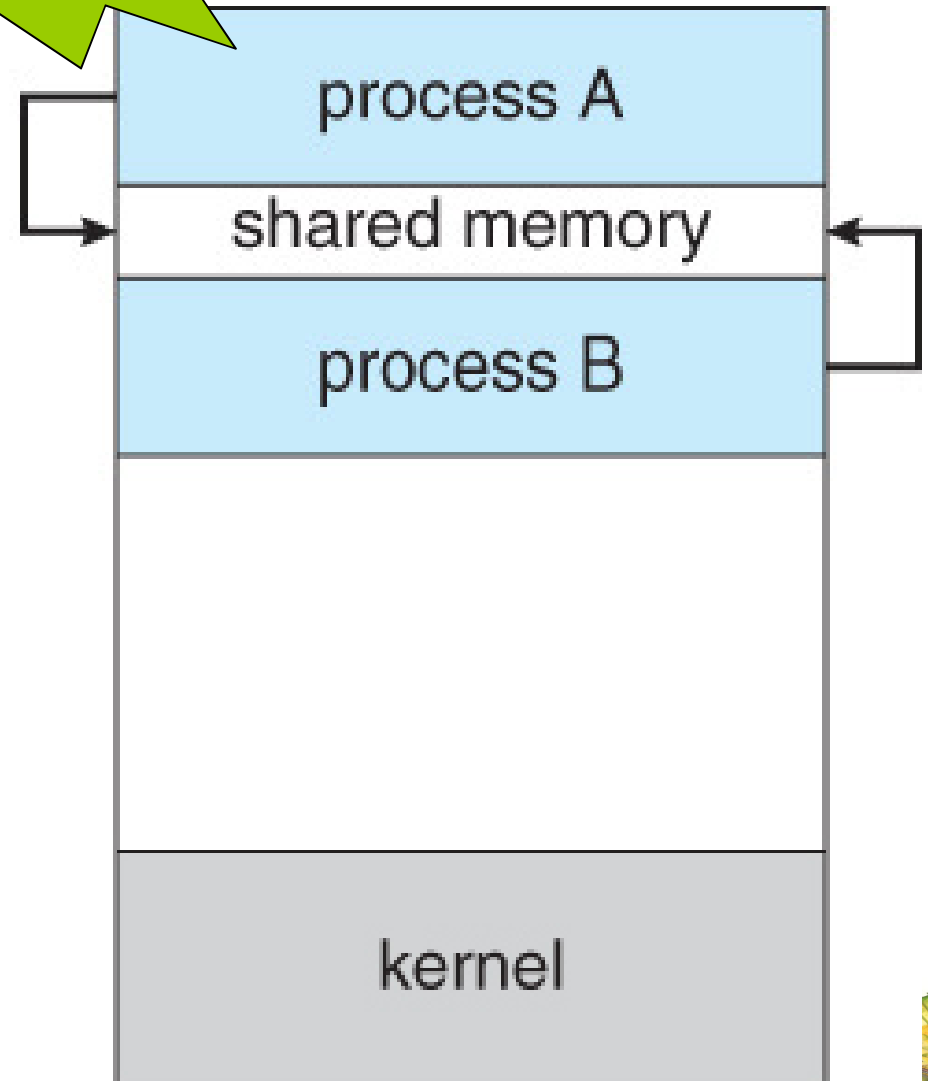


# Message Passing (Send, Receive)

Communicator(Socket)  
Host IP, port/process #



(a)



(b)





# Shared-Memory Issues

How one process can pass information to another process

Emerging critical activities: two or more process do not get into each other's way

Process dependency needs proper sequencing

How about Threads??





# Producer-Consumer Problem

- Problem: There is a set of resource buffers shared by producer and consumer threads
- **Producer** inserts resources into the buffer set
  - ◆ Output, disk blocks, memory pages, processes, etc.
- **Consumer** removes resources from the buffer set
  - ◆ Whatever is generated by the producer
- Producer and consumer execute at different rates
  - ◆ No serialization of one behind the other
  - ◆ Tasks are independent (easier to think about)
  - ◆ The buffer set allows each to run without explicit handoff
- Paradigm for cooperating processes,
  - **unbounded-buffer** : places no practical limit on the size of the buffer
  - **bounded-buffer** : assumes that there is a fixed buffer size






# Bounded-Buffer – Shared-Memory Solution

## ■ Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```





# bounded buffer

## PRODUCER

```
while (true) {  
    /* produce an item and put in  
       nextProduced */  
    while (count == BUFFER_SIZE)  
        do nothing;  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

## CONSUMER

```
while (true) {  
    while (count == 0)  
        do nothing  
    nextConsumed = buffer [out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
    /* consume the item in  
       nextConsumed */  
}
```





## Count=5 @ shared memory

| Producer                                    | Consumer                                    |
|---|---|
| Count++                                     | Count--                                     |
| 4, 5, 6 !                                   |   |
| Reg1=Counter<br>Reg1=Reg1+1<br>Counter=Reg1 | Reg2=Counter<br>Reg2=Reg2-1<br>Counter=Reg2 |
| However, following situation may arise:     |   |
| Reg1=counter      5                         |   |
| Reg1=Reg1+1      6                          | Reg2=counter      5                         |
|   | Reg2=Reg2-1      4                          |
| Counter=Reg1                                | Counter=Reg2                                |
| Race Condition                              |   |

Nondeterministic  
Behavior  
**4/6 !**

Several processes access and manipulate the same data concurrently and the outcome of the execution depends on the particular order in which the access takes place.






# How to avoid Race Condition?

---

- Prohibit more than one process from reading and writing the shared data at the same time.
  
- **Mutual exclusion:**
  - making sure that if one process is using a **shared variable/file**, the other processes will be excluded from doing the same.
  
- **Critical section/Critical region:**
  - part of the program where the shared memory is accessed.





do{

entry section

**Critical Section** (updating variables, table, file etc)

exit section

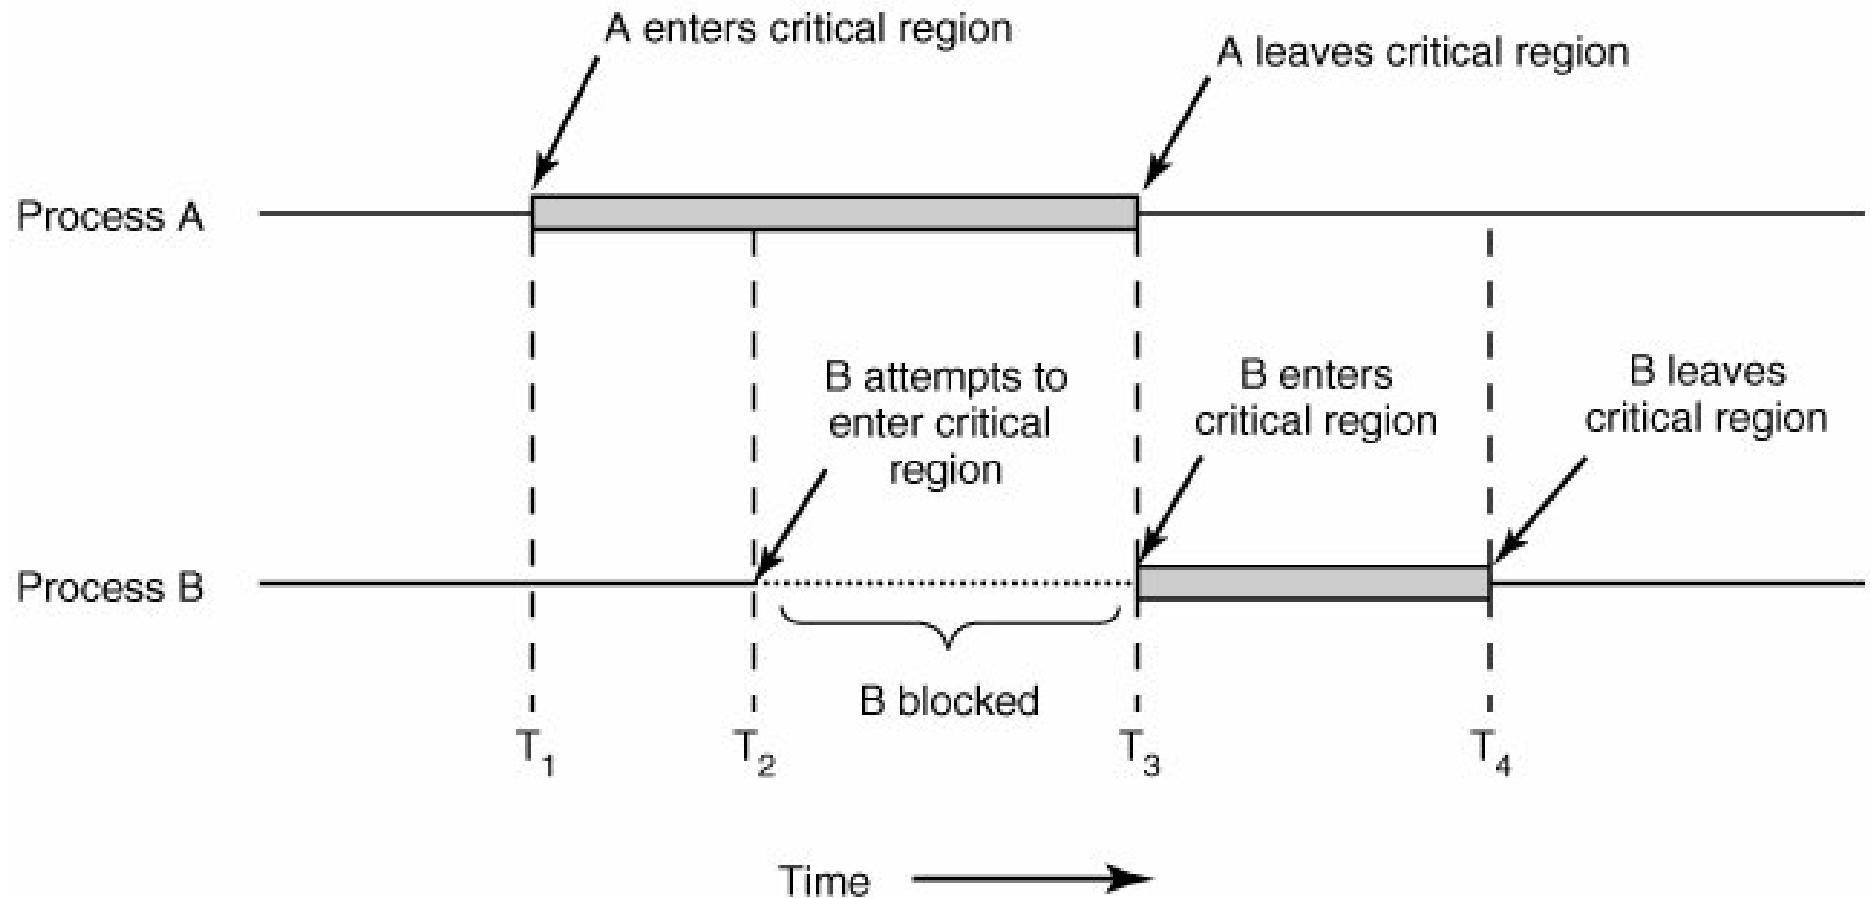
**remainder section**

} while(1);





# Mutual Exclusion using critical regions





# Critical Section

## Four Conditions to have a good solution

---

### Mutual Exclusion

1. No two processes may be simultaneously inside their critical regions.

### Relative Speed

2. No assumptions may be made about speeds or the number of CPUs.

### Progress

3. No process running outside its critical region may block other processes.

### Bounded Wait

4. No process should have to wait forever to enter its critical region.

Bounded # of time other process can enter the critical section, between a process request to a process granted to enter the critical section





# Reason Of Race Condition

---

- Context Switch

- allow concurrent (seems) access to the critical section.

- Sol<sup>n</sup>:

Why not stop context switching during critical section execution.





# Disable Interrupt

## — Disabling Interrupts

```
{  
    ...  
    disableINT();  
    critical_code();  
    enableINT();  
    ...  
}
```

## Problems:

- ▶ It's not wise to give user process the power of turning off INTs.
  - ▶ Suppose one did it, and never turned them on again
- ▶ useless for multiprocessor system **Disable/Enable message needs to pass other processors. Takes time and decrease performance**

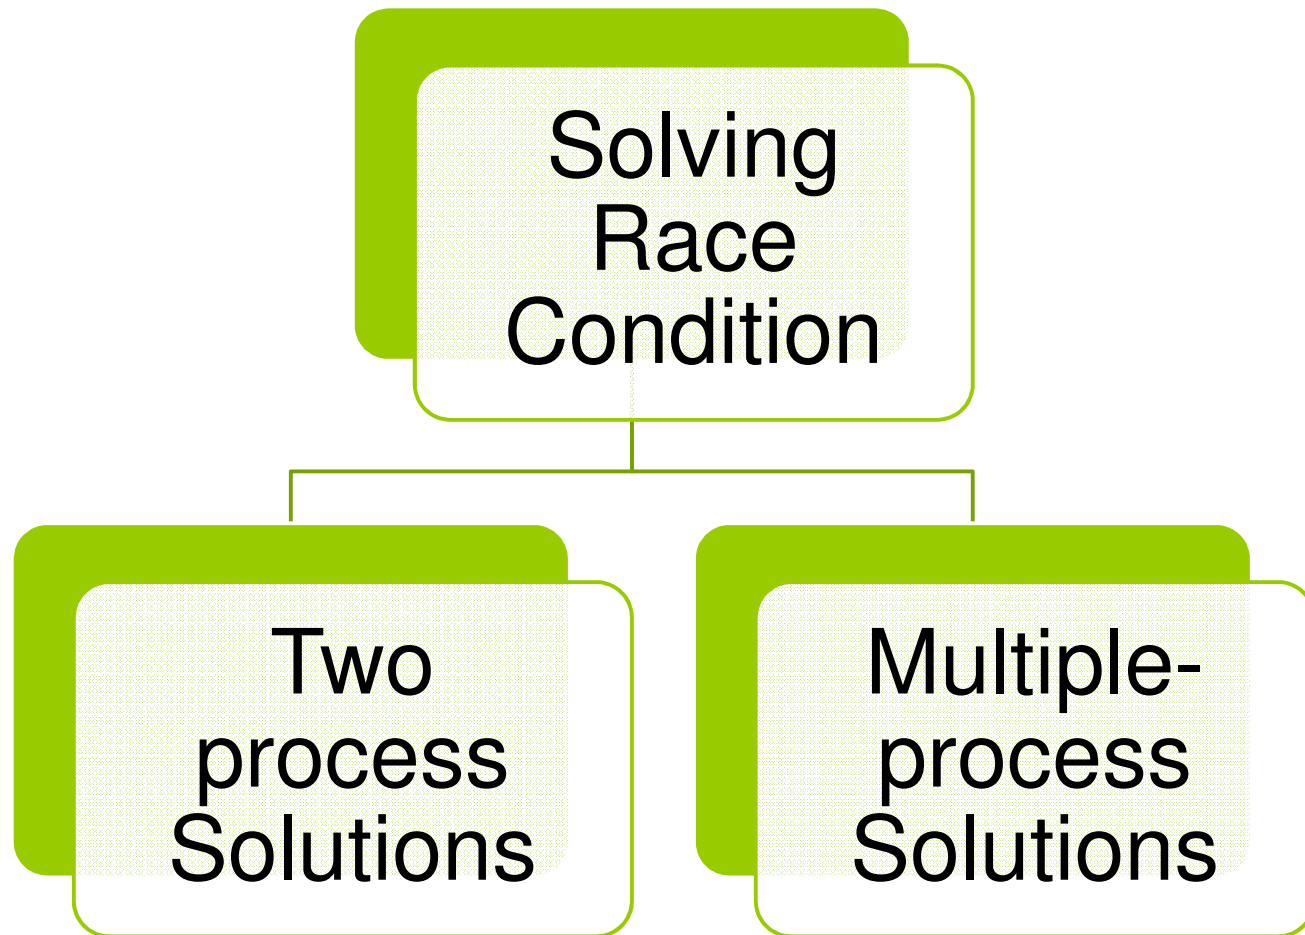
Disabling INTs is often a useful technique within the kernel itself but is not a general mutual exclusion mechanism for user processes.








# Solving Race Condition





# Mutual Exclusion with Busy Waiting- Lock Variable Method

## — Lock Variables



```
1 int lock=0; //shared variable
2 {
3     ...
4     while(lock); //busy waiting
5     lock=1;
6     critical_code();
7     lock=0;
8     ...
9 }
```

**Violates  
Condition ME**

### Problem:

- ▶ What if an interrupt occurs right at line 5?
- ▶ Checking the lock again while backing from an interrupt?





# Algorithm 1: Strict Alternation

— Strict Alternation

initially turn=0

Spin Lock

## Process 0

turn==1

```
1 while(TRUE){
2   while(turn != 0);
3   critical_region();
4   turn = 1;
5   noncritical_region();
6 }
```

## Process 1

```
1 while(TRUE){
2   while(turn != 1);
3   critical_region();
4   turn = 0;
5   noncritical_region();
6 }
```

turn==0

**Problem: violates Relative Speed/ Progress as well**

- ▶ One process can be blocked by another not in its critical region.
- ▶ Requires the two processes strictly alternate in entering their critical region.





## Algorithm 2: Modified Strict Alternation

**Problem Strict Alternation:** Not sufficient information about process state.  
Only remember-which process is allowed to enter Critical Section.  
"turn = 0/1"

**Solution:** Replace "turn" with a Boolean array flag[2];  
flag[0]=flag[1]=0;

### Process 0

```
while(TRUE)
{
    flag[0]=1;
    while (flag[1]);

    Critical Section

    flag[0]=0;
}
```

Context Switch

### Process 1

```
while(TRUE)
{
    flag[1]=1;
    while (flag[0]);

    Critical Section

    flag[1]=0;
}
```

**Mutual Exclusion-OK**

**Relative Speed- No problem**

**Progress - Not OK**

- start same time  
context switch problem





# Switching Order : while & flag

## Results two (2) processes @ CS

flag[0]=flag[1]=0;

### Process 0

```
while(TRUE)
{
    while (flag[1]);
    flag[0]=1;

    Critical Section

    flag[0]=0;
}
```

### Process 1

```
while(TRUE)
{
    while (flag[0]);
    flag[1]=1;

    Critical Section

    flag[1]=0;
}
```





# Algorithm 3: Peterson's Solution

— Peterson's Solution

Combined the idea of  
-lock variable  
-warning variable

```
int interest[0] = 0;  
int interest[1] = 0;  
int turn;      turn=0;
```

**P0**

```
1 interest[0] = 1;  
2 turn = 1;  
3 while(interest[1] == 1  
4       && turn == 1);  
5 critical_section();  
6 interest[0] = 0;
```

**P1**

```
1 interest[1] = 1;  
2 turn = 0;  
3 while(interest[0] == 1  
4       && turn == 0);  
5 critical_section();  
6 interest[1] = 0;
```



Wikipedia – Peterson's algorithm


[http://en.wikipedia.org/wiki/Peterson%27s\\_algorithm](http://en.wikipedia.org/wiki/Peterson%27s_algorithm)





# Multiple-Process Solution

## Bakery Algorithm



Bakeries, Ice-cream Shops



Lowest Token is served next

Distributed Deterministic

Two Customers can receive same numbers

### Need a Tie Breaker

```
do {
    choosing[i] = true;
    number[i] = max(number[0], number[1], ..., number [n - 1])+1;
    choosing[i] = false;
    for (j = 0; j < n; j++) {
        while (choosing[j]) ;
        while ((number[j] != 0) && (number[j,j] < number[i,i])) ;
    }
    critical section
    number[i] = 0;
    remainder section
} while (1);
```

|             |             |
|-------------|-------------|
| number[j],j | number[i],i |
|-------------|-------------|



- Notation  $\leq$  lexicographical order (ticket #, process id #)
  - ◆  $(a,b) < c,d$  if  $a < c$  or if  $a = c$  and  $b < d$
  - ◆  $\max(a_0, \dots, a_{n-1})$  is a number,  $k$ , such that  $k \geq a_i$  for  $i = 0, \dots, n-1$

- Shared data

**boolean** choosing[n];

**int** number[n];

Data structures are initialized to **false** and **0** respectively





# Example: Bakery Algorithm

|       | Process 0   | Process 1   | Process 2   |
|-------|-------------|-------------|-------------|
|       | do{         | do{         | do{         |
| Case1 | number[0]=1 | number[1]=2 | number[2]=3 |
| Case2 | number[0]=1 | number[1]=2 | number[2]=1 |
|       | } while(1); | }while(1);  | }while(1);  |

**Case1:** all processes arrive in order.

$(1, 0 < 2, 1), (2, 1 < 3, 2)$

0 in CS      1 in CS

**Case2:** all processes arrive same time.

$(1, 0 < 1, 2)$

$(2, 1 < 1, 2)$  [process 2 arrives before 1]

0 in CS

2 in CS

Always lowest token  
will be served next

```
do {
    choosing[i] = true;
    number[i] = max(number[0], number[1], ..., number[n - 1]) + 1;
    choosing[i] = false;
    for (j = 0; j < n; j++) {
        while (choosing[j]);
        while ((number[j] != 0) && (number[j, j] < number[i, i]));
    }
    critical section
    number[i] = 0;
    remainder section
} while (1);
```







# Hardware Solution: TSL instruction

## ■ Atomic Operation:

some operations those **read and change data** within a **single, uninterruptible** step.

### **Example: Test and Set Lock instruction**

TSL returns the current value of a memory location and replaces it with a given new value- TRUE

**Initially Lock = 0**

```
boolean TestAndSet ( boolean & target)
{
    boolean rv= target;    Two memory cycles
    target= true;          1) Read
    return rv;             2) Write
}
```

**Why does it call- H/W Solution??**

```
do
{
    while(TestAndSet(Lock));

    Critical Section

    lock=false;

}while(1);
```



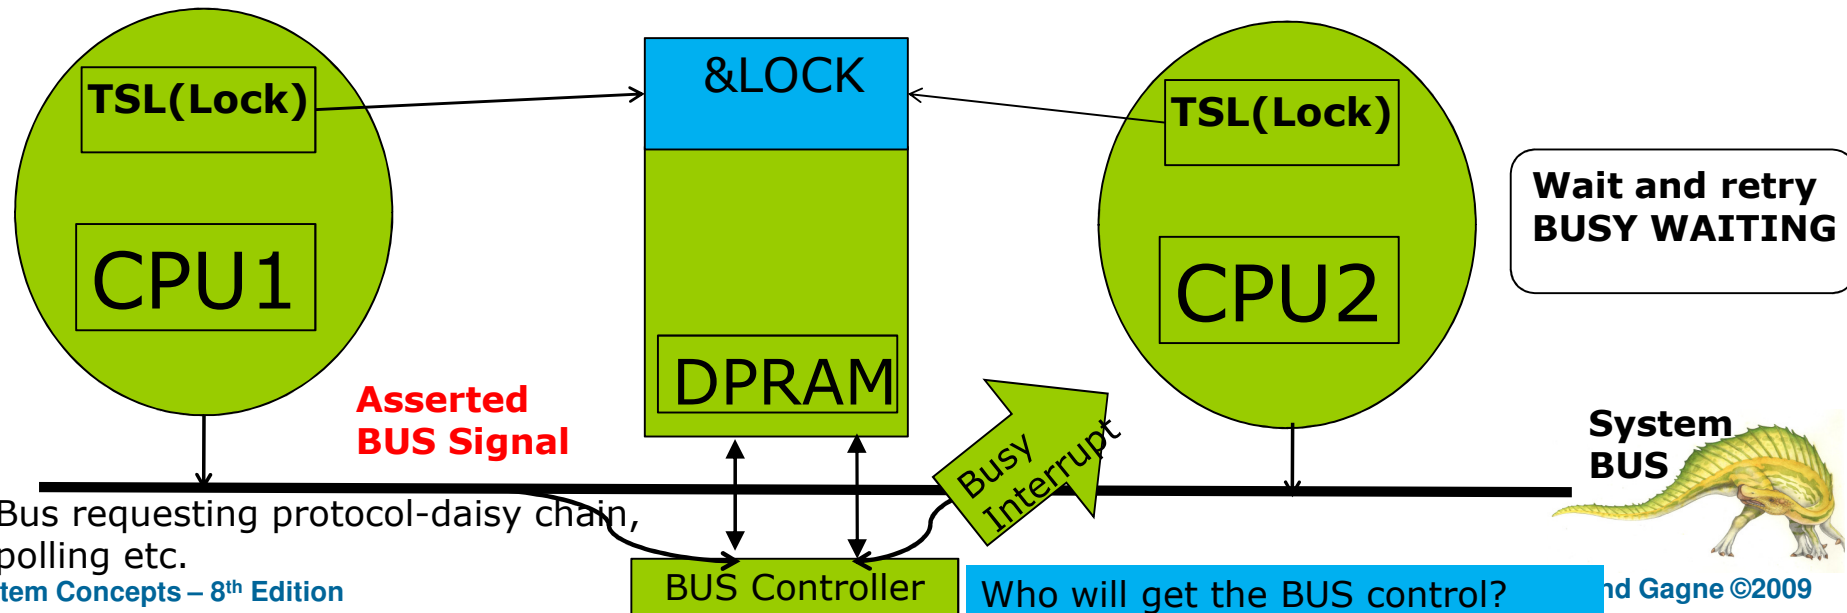
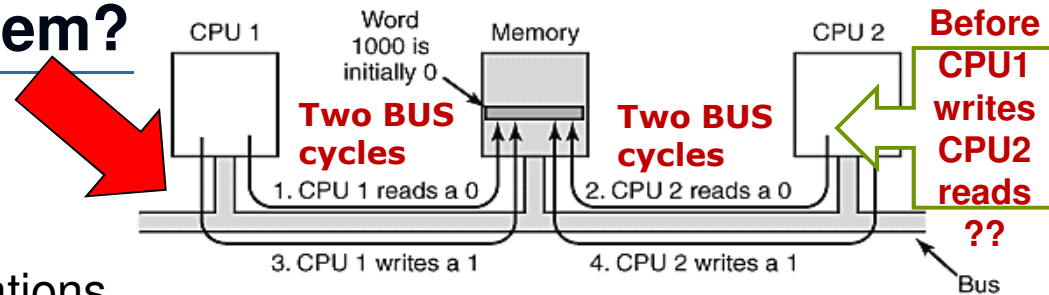
# Why does it call- Hardware Solution?

## How about Multiprocessors system?

### Solution:

#### Hardware Lock (Memory Barrier):

- Provide synchronization on memory operations
- A processor generates signal and that helps to prevent other processors from using the system bus.
- Example: Lock instruction 8086(intel)-a full barrier





- Both Peterson's solution and solution using TSL are correct
  - but both have BUSY WAITING.
  - Wasting CPU cycle.





# Sleep and Wakeup System Calls

Sleep & Wakeup • • •

Solves Busy Waiting

```
#define N 100  /* number of slots in the buffer */  
int count = 0; /* number of items in the buffer */
```

```
1 void producer(){  
2   int item;  
3   while(TRUE){  
4     item = produce_item();  
5     if(count == N)  
6       sleep();  
7     insert_item(item);  
8     count++;  
9     if(count == 1)  
10      wakeup(consumer);  
11   }  
12 }
```

```
1 void consumer(){  
2   int item;  
3   while(TRUE){  
4     if(count == 0)  
5       sleep();  
6     item = rm_item();  
7     count--;  
8     if(count == N - 1)  
9       wakeup(producer);  
10    consume_item(item);  
11  }  
12 }
```





# Sleep and Wakeup System Calls

## Sleep & Wakeup

```
#define N 100 /* number of slots in the buffer */  
int count = 0; /* number of items in the buffer */
```

```
1 void producer(){  
2     int item;  
3     while(TRUE){  
4         item = produce_item();  
5         if(count == N)  
6             sleep();  
7         insert_item(item);  
8         count++;  
9         if(count == 1)  
10            wakeup(consumer);  
11     }  
12 }
```

```
1 void consumer(){  
2     int item;  
3     while(TRUE){  
4         if(count == 0)  
5             sleep();  
6         item = rm_item();  
7         count--;  
8         if(count == N - 1)  
9             wakeup(producer);  
10        consume_item(item);  
11    }  
12 }
```

Producer and  
consumer

Both can be  
in sleep()  
forever

Deadlock!





# Producer-Consumer Problem

— Race Condition

## Problem

1. Consumer is going to sleep upon seeing an empty buffer, but INT occurs;
2. Producer inserts an item, increasing `count` to `1`, then call `wakeup(consumer)`;
3. But the consumer is not asleep, though `count` was `0`. So *the `wakeup()` signal is lost*;
4. Consumer is back from INT remembering `count` is `0`, and goes to sleep;
5. Producer sooner or later will fill up the buffer and also goes to sleep;
6. Both will sleep forever, and waiting to be waken up by the other process. Deadlock!





# Producer-Consumer Problem

— Race Condition

Solution: Add a *wakeup waiting bit*

1. The bit is set, when a wakeup is sent to an awoken process;
2. Later, when the process wants to sleep, it checks the bit first. Turns it off if it's set, and stays awake.

What if many processes try going to sleep?





# Higher-Level Synchronization

---

- We looked at using locks to provide mutual exclusion
- Locks work, but they have some drawbacks when critical sections are long
  - Spinlocks – inefficient
  - Disabling interrupts – can miss or delay important events
- Instead, we want synchronization mechanisms that
  - Block waiters
  - Leave interrupts enabled inside the critical section
- Look at two common high-level mechanisms
  - **Semaphores**: binary and counting semaphore
  - **Monitors**: mutexes and condition variables







# What is a Semaphore?

- ▶ A locking mechanism
  - ▶ An integer or ADT
- that can only be operated with:
- Block waiters, interrupts enabled within CS**  
**Hide implementation view.**

**During Semaphore updating:  
No other process can access  
the Semaphore**

**Automatically  
done all process  
by a single  
command**

## Atomic Operations

|             |             |
|-------------|-------------|
| P()         | V()         |
| Wait()      | Signal()    |
| Down()      | Up()        |
| Decrement() | Increment() |
| ...         | ...         |

```
down(S) {  
    while(S <= 0);  
    S--;  
}
```

```
up(S) {  
    S++;  
}
```

**Randomly Wakeup  
One process**

More meaningful names:

- ▶ increment\_and\_wake\_a\_waiting\_process\_if\_any()
- ▶ decrement\_and\_block\_if\_the\_result\_is\_negative()





# Semaphore

---

## How to ensure atomic?

1. For single CPU, implement `up()` and `down()` as system calls, with the OS disabling all interrupts while accessing the semaphore;
2. For multiple CPUs, to make sure only one CPU at a time examines the semaphore, a lock variable should be used with the TSL instructions.

*This section is based on (Tanenbaum, 2001)*





# Semaphore is a Special Integer

A semaphore is like an integer, with three differences:

1. You can initialize its value to any integer, but after that the only operations you are allowed to perform are *increment* ( $S++$ ) and *decrement* ( $S--$ ).
2. When a thread decrements the semaphore, if the result is negative ( $S \leq 0$ ), the thread blocks itself and cannot continue until another thread increments the semaphore.
3. When a thread increments the semaphore, if there are other threads waiting, one of the waiting threads gets unblocked.

*This section is based on (Tanenbaum, 2001)*





# Still busy waiting? Sol- process block

**wait(s):**

```
s.count=s.count-1;  
if(s.count < 0) then  
{  
    place this process in s.queue;  
    block this process  
}
```

```
typedef struct {  
    int value;  
    struct process *queue;  
} semaphore;
```


**signal(s):**

```
s.count=s.count+1;  
  
if(s.count <= 0) then  
{  
    remove a process P from s.queue;  
    place process P on ready queue  
}
```





# Semaphore Types





# Binary Semaphore

## ADT- Hide details implementation from user

wait( )

{

if (BS. value=1) then  
    BS.value=0

else

    Place process in S.Queue and  
    block the process

}

signal( )

{

if (BS. Queue is empty)  
    BS.value=1

else

    Remove a process from BS.Queue  
    and place p in ready queue.

}

**WHY: do we need Binary semaphore code?**

**Can't we use COUNTING semaphore with value 1?**

**Answer:** They are different, BS supports value 1/0.

Where as: **signal()** before **wait()** makes semaphore value to more than 1.







# Example


**i++ is not atomic in assembly language**

```
LOAD    r0, [i]    ;load the value of 'i' into
                   ;a register from memory
ADD     r0, r0, 1    ;increment the value
                   ;in the register
STORE   r0, [i]     ;write the updated
                   ;value back to memory
```

Interrupts might occur in between. So, **i++** needs to be protected.

*This section is based on (Tanenbaum, 2001), Pages 103 – 105  
slides taken from youtube*





# Example: Binary and Counting Semaphore

| Binary   |  | Counting   |                                  |
|--|--|--|----------------------------------|
| BS=1<br>T1, T2                                   |  | S=3<br>T1, T2, T3, T4, T5                        |                                  |
| wait()<br>i++ ; critical region (CR)<br>Signal() |  | wait()<br>i++ ; critical region (CR)<br>Signal() |                                  |
| STEP 1   | T1 enters CR , BS=0                    |  | T1 enters CR, S=2                |
| STEP 2   | T2 want to enters CR<br>BS.Queue[0]=T2 |  | T2 enters CR, S=1                |
| STEP 3   | T1 exits CR , BS=1, remove T2          | STEP 3   | T3 enters CR, S=0                |
| STEP 4   | T2 enters CR , BS=0                    | STEP 4   | T4 enters CR, S=-1 S.Queue[0]=T4 |
|  |  | STEP 5   | T5 enters CR, S=-2 S.Queue[1]=T5 |
|  |  | STEP 6   | T1 exits CR S=-1 T5 enters       |
|  |  | STEP 7   | T2 exits CR S =0 T4 enters       |
|  |  |  |                                  |







# Mutex: A Special BS

## mutex\_lock:

```
TSL REGISTER, MUTEX  
CMP REGISTER, #0  
JZE ok_go_CR
```

```
JMP mutex_lock
```

**ok\_go\_CR :**  
critical section entered

## mutex\_unlock:

```
MOVE MUTEX, #0
```

```
RET to caller
```

| BS  | MUTEX                               |
|---|-------------------------------------|
| P/T can be lock itself and unlock by others | P/T locks itself and unlocks itself |
| Work with signal                            | Work with Token                     |
| Traffic Light                               | ATM Booth                           |
| 1 to enter CR                               | 0 to enter CR                       |





# Problem with Semaphores

- Programmers must have knowledge to use semaphores.
- Otherwise: deadlock situation may arrive.

```
init(&sem1, 0);
```

```
init(&sem2, 0);
```

## Process One

```
a1;  
wait(&sem2);  
signal(&sem1);  
a2;
```

## Process Two

```
b1;  
wait(&sem1);  
signal(&sem2);  
b2;
```

# Deadlock





# Monitor

---

- High level synchronization construct
  - defined with collection of procedures, variables and data structures
- Only one process can active in a monitor at a time.
- Up to the compiler (trace monitor specially)
  - how to implement **Mutual Exclusion (ME)**
  - less likely to get something wrong
- Programmer just select critical region and put code inside monitor
  - compiler determines the rest

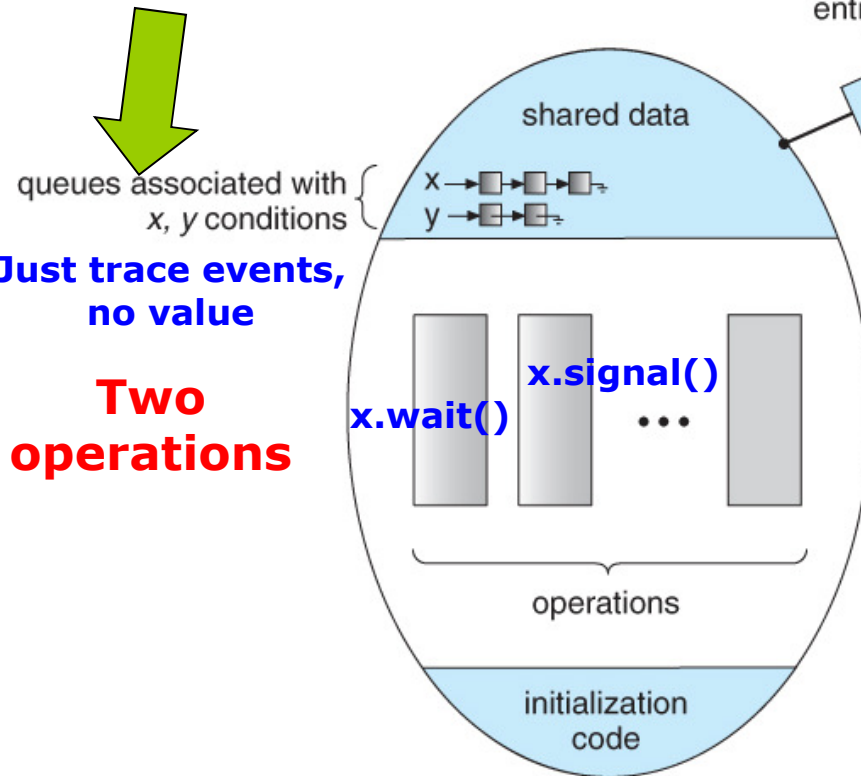




# Monitor: Implementation

```
monitor monitor-name {  
    shared variable declarations  
  
    procedure body P1 ( . . . )  
        . . .  
}  
  
    procedure body P2 ( . . . )  
        . . .  
}  
.  
.  
.  
{  
    initialization code  
}  
}
```

1. One process at a time can be active in monitor.
2. A thread/process @ monitor may have to block itself because of its request may not complete immediately.



**Just trace events,  
no value**

**Two  
operations**





# Monitor Example

Synchronization mechanisms need more than just mutual exclusion; also need a way to wait for another thread/process to do something (e.g., wait for a character to be added to the buffer)

Process needs to block.  
How they know buffer is full or empty?

Conditional variable with wait() & signal()

## Step 1

C P m

Entry Queue

full empty



## Step 2

P C m

Entry Queue

full empty



```

1 monitor ProducerConsumer
2   condition full, empty;
3   integer count;
4
5   procedure insert(item: integer);
6   begin
7       if count = N then wait(full);
8       insert_item(item);
9       count := count + 1;
10      if count = 1 then signal(empty)
11      end;
12
13  function remove: integer;
14  begin
15      if count = 0 then wait(empty);
16      remove = remove_item;
17      count := count - 1;
18      if count = N - 1 then signal(full)
19      end;
20      count := 0;
21  end monitor;
```

```

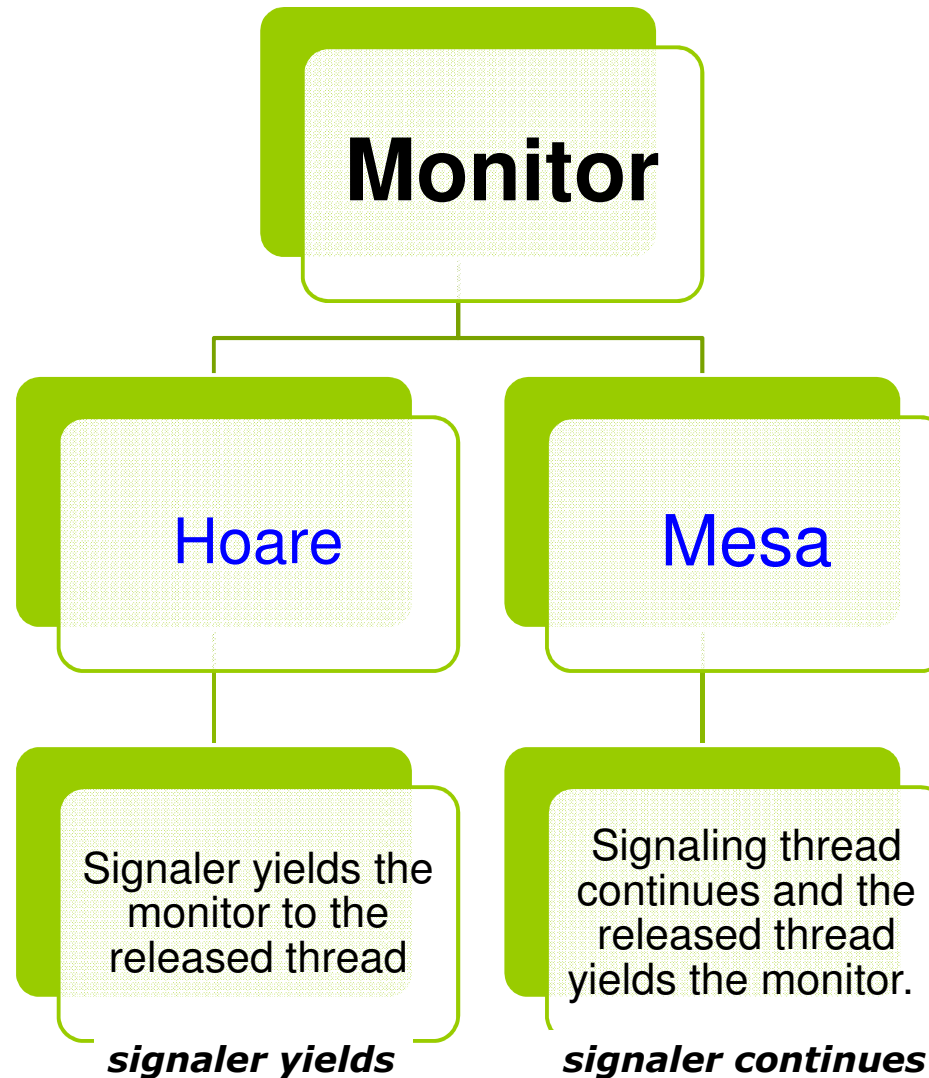
1 procedure producer;
2 begin
3     while true do
4     begin
5         item = produce_item;
6         ProducerConsumer.insert(item)
7     end
8 end;
9
10 procedure consumer;
11 begin
12     while true do
13     begin
14         item = ProducerConsumer.remove;
15         consume_item(item)
16     end
17 end;
```






# Who will hold the Monitor Lock!

Signaling process / Released process



# Reading 7.1-7.4, 7.6-7.7

---





# Race Conditions

- Multiple threads/processes attempt to access a shared resource
  - such as - a shared variable / a shared file,
- at least one of the accesses is an update, and the accesses can result in an error.

